

GAME DAY — April 16, 1pm ET — Laura Small

Four punchlines (sticky note):

1. "I refused to follow when Meta reorged to serve only advertiser needs. Anthropic = mission baked into architecture, not bolted on as policy."
 2. "I'd rather explain a revenue gap than explain why our platform influenced an election."
 3. "You earn trust by solving the problems people already know they have."
 4. "A system working as designed can still be the problem."
-

Q1: TELL ME ABOUT YOURSELF (90 sec)

PURPOSE: What have you done? Consumer PM credentials + bridge to mission.

BEATS:

- HOOK: frontier AI → consumer products → trust is the hard part
- META CONSUMER: Reels, Explore, parental controls, privacy, Instagram/FB/WhatsApp/Messenger, 3B users
- META MISSION: co-founded Responsible AI org, built safety features that cost the company revenue (SEEDS Q2)
- RING: GM, Neighbors safety app, 50M devices
- SOCURE: AI Agent platform, 0-to-1, 300% ARR
- FRACTIONAL CTO: hands-on building with founding teams, not advisory
- CLOSE: technology changes, trust problem doesn't → that's why Anthropic Consumer PM

VERBATIM:

BLUF: "I take frontier AI and turn it into consumer products people trust enough to use every day. I've done it four times at scale, and each time the hardest part wasn't the technology. It was the trust."

FULL: "At Meta I shipped recommendation systems for Reels and Explore for 3 billion users, and built parental controls and privacy infrastructure across Instagram, Facebook, WhatsApp, and Messenger. I also co-founded Meta's Responsible AI org, which meant building consumer safety features that actually cost the company revenue, things like ad fairness constraints and content transparency tools. At Ring I was GM scaling AI features and the Neighbors community safety app across 50 million devices. At Socure I built an AI Agent platform from scratch, grew ARR 300%. In between I've been hands-on as a fractional CTO building products from zero with

founding teams. The thread: the technology changes, the trust problem doesn't. That's what draws me to Consumer PM at Anthropic."

Q2: WHY ANTHROPIC? (75 sec)

PURPOSE: Why here, why now? Picks up where Q1 left off. Don't repeat "I co-founded Responsible AI." Extend.

BEATS:

- "That Responsible AI work" (reference Q1, don't repeat)
- Most meaningful work at Meta, but incentive structure fought us
- Every safety feature reduced a metric someone was measured on
- Anthropic: incentive structure IS the safety work (RSP, interpretability, Pentagon)
- Product side: trust alone doesn't drive daily usage
- CLOSE: build the habit on trust, like Reels/Instagram or Neighbors/Ring, done it 4x
- FALLBACK: if Q1 wasn't asked, open with "I co-founded Meta's Responsible AI org"

VERBATIM:

"That Responsible AI work I mentioned was the most meaningful thing I did at Meta. But the incentive structure worked against us. Safety was always fighting for resources against revenue. We'd ship parental controls on Instagram or fairness constraints on housing ads, and every one reduced a metric someone else was being measured on.

At Anthropic, the incentive structure IS the safety work. The RSP (Responsible Scaling Policy), the interpretability research, the willingness to walk away from a \$200M Pentagon contract. That's not marketing. That's a technical and organizational commitment I haven't seen anywhere else.

On the product side, Claude is the most trusted frontier model, but trust alone doesn't drive daily usage. The consumer product needs someone who can build the habit on top of that trust, the way Reels built a daily habit on top of Instagram's social graph, or the way Neighbors built a daily habit on top of a doorbell camera. I've done that at scale four times."

FALLBACK (if Laura skipped Q1): Replace "That Responsible AI work I mentioned" with "I co-founded Meta's Responsible AI org because I believed it mattered whether we got AI right for the 3 billion people using Instagram, WhatsApp, and Messenger every day. That was the most meaningful work I did there."

Q3: WHY ANTHROPIC FISHING FOLLOW-UP (30 sec)

BEATS:

- Three years thinking about responsible AI at a company that actually wants it
- Structurally true, not aspirationally true
- IF PUSHED: hearing aids since 2, technology made me equal, career = making that true for more people

VERBATIM:

"Honestly? I've spent three years since Meta thinking about what it would look like to do responsible AI work at a company where the company actually wants you to do it. Where you're not fighting the org chart to do the right thing. Anthropic is the only place where that feels structurally true, not just aspirationally true."

If pushed further: "I've worn hearing aids since I was two. Technology is the reason I can have this conversation. Every role I've taken has been me trying to make that true for more people. Anthropic is where that work matters most right now."

Q4: AI SAFETY (90 sec)

BEATS:

- Tension: what makes AI great (seamless, invisible) = what makes failures harder to catch
- RING NEIGHBORS: motion alerts too trusted, nobody caught errors
- META REELS/EXPLORE: rec systems, feedback loops too slow at scale
- SOCURE: identity errors = can't open bank account, insisted human-in-the-loop
- PATTERN: more capable → more trust → less checking → higher stakes
- CLOSE: safety embedded in product decisions, not separate workstream

VERBATIM:

"One tension I think about a lot from my time at Meta: the things that make AI products great for users, seamless integration, low friction, feeling invisible, are exactly what make safety failures harder to catch. When the tool disappears into the workflow, users stop double-checking."

At Ring we saw this with the Neighbors app. AI-generated motion alerts were so trusted that when the system misclassified something, nobody caught it until the damage was done. At Meta we saw it with Reels and Explore recommendation systems surfacing content at scale where the feedback loops were too slow to catch emerging harms. And at Socure, identity verification

errors don't just mean a bad user experience. They mean someone can't open a bank account or rent an apartment. We insisted on human-in-the-loop for high-risk decisions when the business wanted full automation.

The pattern across all three: the more capable and trusted the AI system becomes, the higher the stakes when it fails, and the less likely anyone is to notice in time. That's why I think safety can't be a separate workstream from product. It has to be embedded in the product decisions themselves. How you design feedback loops, how you set defaults, how you surface uncertainty to users. Those are product choices with safety consequences.

The question isn't whether to ship. It's how you build the feedback loops that catch problems before users do. That's a product design question as much as a research question."

Q5: CONSUMER PRODUCT YOU'VE BUILT (2 min)

BEATS:

- Meta parental controls across Instagram, Facebook, Messenger
- Tension: visible for parents, not so restrictive teens make burners or lie about age
- Tested 3 models: opt-in, default-on, supervised
- Shipped supervised: transparency + autonomy, 35% higher adoption
- Refused default-on (faster metrics, drives circumvention)
- CLOSE: chose trust over speed
- CLAUDE BRIDGE: family account, teen + mental health, same tension

VERBATIM:

"At Meta, I led the team that redesigned the parental controls and privacy experience across Instagram, Facebook, and Messenger. This wasn't a settings page buried in a menu. We were building the product that parents would use to decide whether their kid could use Instagram safely.

The hardest part wasn't the engineering. It was the tension between making the controls visible enough that parents trusted them, but not so restrictive that teenagers just created burner accounts or lied about their age to bypass them entirely. We tested three different models: opt-in controls, default-on with opt-out, and a supervised mode that let parents see activity without controlling it.

We shipped the supervised model because the data showed it was the only one where both parents AND teens felt respected. Parents got transparency. Teens got autonomy within guardrails. Adoption was 35% higher than the restrictive model.

What I'm proudest of is that we resisted the pressure to ship the default-on version first, which would have hit our coverage metrics faster but would have pushed teenagers to create secondary accounts or fake their age to get around the controls, undermining the whole system. We chose trust over speed."

Claude bridge if asked: "Imagine a Claude family account where parents want visibility but a teenager needs to feel safe asking about mental health or identity. Same tension. Too restrictive and users route around it. Too permissive and parents don't trust it."

Q6: WHY CONSUMER PM? (60 sec)

BEATS:

- Biggest impact = product millions touch daily
- Research foundational but irrelevant if product doesn't reflect it
- Bridge between research and morning product

VERBATIM:

"Because the biggest impact of safe AI will be in the product that millions of people touch every day. The research is foundational, but it doesn't matter if the consumer product doesn't reflect it. I want to be the person making sure that Anthropic's safety research actually shows up in the product people use every morning. That's the bridge between research and impact, and it's where I've spent my whole career."

Q7: SENIOR DIRECTOR → IC? (60 sec)

BEATS:

- Hands-on last 3 years (Socure: specs, customer calls, UX. Fractional CTO: building from zero)
- Yes, 200+ teams too, but the work I keep choosing is building
- Roles change every 3-6 months at Anthropic
- CLOSE: smoke jumper, not fixed org chart

VERBATIM:

"I've been hands-on for the last three years straight. At Socure I built the AI Agent platform from scratch. Writing product specs, sitting in customer calls, iterating on the UX weekly. As a fractional CTO/CPO I was embedded with founding teams, building products from zero, not

advising from a distance. Yes, I've also done the big org leadership job, built teams of 200+, but the work I keep choosing is the building. At Anthropic specifically, the roles change fast. What you start doing could change in 3-6 months. I'd rather be a smoke jumper who goes where the fire is than a director who owns a fixed org chart."

Q8: WHAT SINCE META? (45 sec)

BEATS:

- Three things, all hands-on
- VC: 120+ startups, 3x returns, pattern recognition
- Fractional CTO: embedded with founders, consumer + enterprise from zero
- Socure: AML business, AI Agent platform, 300% ARR
- Foundation: Meta consumer (Reels, Explore, Instagram, WhatsApp, parental controls)
- CLOSE: concentrate on one product, Anthropic

VERBATIM:

"Three things, all hands-on. Running a small AI-focused VC fund where I evaluated 120+ startups and got roughly 3x returns, which gave me pattern recognition across every category of AI product. Fractional CTO/CPO work where I was embedded with founding teams building consumer and enterprise products from zero. And at Socure, running the AML business and building the AI Agent platform from scratch, growing ARR 300%. I've been building continuously since Meta, just across more surfaces. The consumer work at Meta, Reels, Explore, Instagram privacy, WhatsApp, parental controls, gave me the foundation. Everything since has been testing those instincts in new domains. Now I want to concentrate on one product where the work matters. That's Anthropic."

Q9: WHAT WOULD YOU BUILD? (60 sec)

BEATS:

- Three surfaces (chat, Code, Cowork) converging
- Dispatch + scheduled tasks + Computer Use + plugins = 90% always-on agent
- Consumer PM challenge: non-technical users controlling agent without understanding permissions/tokens
- My approach: progressive autonomy, Claude earns trust through small tasks first
- CLOSE: consumer instinct + safety thinking = my career

VERBATIM:

"The product is at an inflection point. Claude has three surfaces now, chat, Code, and Cowork, and they're converging. Dispatch creates a persistent thread across phone and desktop. Scheduled tasks add time triggers. Computer Use adds action capability. The plugin architecture adds domain expertise. Stack those together and you're 90% of the way to an always-on agent platform.

The consumer PM challenge is: how do you give non-technical users enough control over an always-on agent that they trust it, without requiring them to understand permission models or token economics? That's a UX problem, a safety problem, and a business model problem all at once. Progressive autonomy, where Claude earns trust through small tasks before being handed bigger ones, is how I'd approach it.

That's the intersection of consumer product instinct and safety thinking, and it's exactly where I've spent my career."

STOCK TO ZERO / 180?

BEATS:

- Hard to answer before it's real
- PBC, Pentagon, Constitutional AI = not reversible in a board meeting
- Already lived the reverse (Meta reorg)
- Pushed back and left. Not hypothetical.

VERBATIM:

"That's hard to answer honestly before it's real. But the reason I'm drawn to Anthropic over competitors is specifically because the mission is structural. PBC (Public Benefit Corporation) status. The Pentagon decision. Constitutional AI baked into training. Those aren't things you walk back in a board meeting. And I've already been through the version where the company reverses course. I know what I did. I pushed back and left."

STORIES (if culture topics come up early)

M1: META RESPONSIBLE AI

- Rec + ad systems optimizing without guardrails
- Built: transparency, regulated ad constraints, parental controls (burners + age lies)
- Every feature reduced revenue

- Reorg came, told to get in line, refused to follow
- FOR LAURA: "I chose not to follow." (agency)
- FOR CULTURE INTERVIEW LATER: "It cost me my seat." (cost of conviction)
- PUNCHLINE (Laura): I refused to follow when Meta reorged. Anthropic = mission baked into architecture.

R2: RING NEIGHBORS (I was wrong)

- Engagement strong but fear-based, some racially biased
- Numbers blinded me, conflated "engaged" with "safe"
- Changed metric to safety perception, invested in moderation
- PUNCHLINE: system working as designed can still be the problem

S1: SOCURE ENTRY

- New GM, founder-led, team skeptical
- 30 days listening, built around THEIR diagnosis
- Months → weeks onboarding, 40% YoY, 300% ARR
- PUNCHLINE: earn trust by solving problems people already know they have

P1: PERSONAL (deeper layer only if asked)

- Hearing aids since 2, technology = why I'm here
- Parents: grad degrees unrecognized, rebuilt from scratch, Queens basement
- Wall Street hollow, Ring parents' peace of mind, Meta stakes real for 3B
- Sister left Lockheed (missiles). Family value system.

USER RESEARCH (real, from power user call April 14 — use in Q9 or product discussion)

- "Consumers don't think in files and folders." Personal knowledge base, auto-organized.
- Connectors broken for personal use: Google Drive root only, Slack no notifications, one Google account limit
- Sharing = "pods of conversations." Distill, fork, share selectively. Not group chat.
- Family plan like Spotify: individual privacy, selective sharing, contextual suggestions

- Health management = 3 products: episodic (knee surgery), trending (15yr blood work), coaching (retirement fear)
 - Cold start: cohort data → individual. Let users accelerate by sharing existing feeds.
 - KEY QUOTE: "It should know better than me what that thing is. But shouldn't act without my approval."
 - DEPLOY AS: "I've been doing user research with power users. Biggest gap isn't capability, it's that the product thinks in files/folders and consumers don't."
-

REMINDERS:

- Flip tempo at minute 15. Ask questions back.
- Don't quote the values page. Ever.
- If you cite any Anthropic decision, go two levels deep on WHY.
- Be a person, not a candidate.